

Logistic Regression Homework

Load the faraway package and then type “wbca” at the prompt... but without the quotes. Read about this dataset here: <https://cran.r-project.org/web/packages/faraway/faraway.pdf>. There are 681 cases of potentially cancerous tumors of which 238 are actually malignant. Determining whether a tumor is really malignant is traditionally determined by an invasive surgical procedure. The purpose of this study was to determine whether a new procedure called fine needle aspiration which draws only a small sample of tissue could be effective in determining tumor status.

- (a) Fit a binomial regression with Class as the response and the other nine variables as predictors. Just do

```
model1 <- glm(Class ~ ., data = wbca, family=binomial)
summary(model1)
```

Report the residual deviance and associated degrees of freedom. Can this information be used to determine if this model fits the data well (think in terms of sample sizes)? If so, what does the chi-squared test for the residual deviance indicate?

Here's the output:

Call:

```
glm(formula = Class ~ ., family = binomial, data = wbca)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.48282	-0.01179	0.04739	0.09678	3.06425

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	11.16678	1.41491	7.892	2.97e-15 ***
Adhes	-0.39681	0.13384	-2.965	0.00303 **
BNucl	-0.41478	0.10230	-4.055	5.02e-05 ***
Chrom	-0.56456	0.18728	-3.014	0.00257 **
Epith	-0.06440	0.16595	-0.388	0.69795
Mitos	-0.65713	0.36764	-1.787	0.07387 .
NNucl	-0.28659	0.12620	-2.271	0.02315 *
Thick	-0.62675	0.15890	-3.944	8.01e-05 ***
UShap	-0.28011	0.25235	-1.110	0.26699
USize	0.05718	0.23271	0.246	0.80589

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 881.388 on 680 degrees of freedom
Residual deviance: 89.464 on 671 degrees of freedom
AIC: 109.46

Number of Fisher Scoring iterations: 8

In order that the Chi-square approximation to work well, we need at least several observations of each treatment... for example, we would need 5 or ten observations where Adhes = 1, BNucl=1, Chrom =3, Epth=2, ... and 5 or ten observations where Adhes=2, BNucl = 1, Chrom = 3, Epith = 2, etc. Probably there are not enough. Even so, I will get the p-value for a chi-squared test:

```
> pchisq(deviance(model1), df.residual(model1), lower=FALSE)
[1] 1
```

This indicates the model fits sufficiently well. This is good, since we're still using all the variables at this point. This is probably one of the better fitting models we could get!

- (b) Use AIC as the criterion in best-subsets method to determine the best subset of variables. This is easy- just do

```
reduced <- step(model1)
summary(reduced)
```

Here's my output:

Call:
glm(formula = Class ~ Adhes + BNucl + Chrom + Mitos + NNucl +
Thick + UShap, family = binomial, data = wbca)

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.44161	-0.01119	0.04962	0.09741	3.08205

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	11.0333	1.3632	8.094	5.79e-16 ***
Adhes	-0.3984	0.1294	-3.080	0.00207 **
BNucl	-0.4192	0.1020	-4.111	3.93e-05 ***
Chrom	-0.5679	0.1840	-3.085	0.00203 **
Mitos	-0.6456	0.3634	-1.777	0.07561 .
NNucl	-0.2915	0.1236	-2.358	0.01837 *
Thick	-0.6216	0.1579	-3.937	8.27e-05 ***
UShap	-0.2541	0.1785	-1.423	0.15461

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 881.388 on 680 degrees of freedom
Residual deviance: 89.662 on 673 degrees of freedom
AIC: 105.66

Number of Fisher Scoring iterations: 8

Again using the Chi-square test:

```
> pchisq(deviance(reduced), df.residual(reduced), lower=FALSE)
[1] 1
```

Checking the null model (model with no predictors, just the intercept term):

```
> pchisq(881, 680, lower=FALSE)
[1] 2.74918e-07
```

- (c) Use the reduced model (the one from part (b)) to predict the outcome (malignant or not) for a patient with

Adhes = 1, BNucl = 1, Chrom = 3, Epith = 2, Mitos = 1,
NNucl = 1, Thick = 4, UShap = 1, USize = 1.

Also, give a 95% confidence interval for your prediction.

```
> nd <- data.frame(Adhes=1, BNucl=1, Chrom = 3, Epith = 2, Mitos = 1, NNucl=1,
Thick=4, UShap=1, USize=1)
> predict(model1, nd, type="response")
1
0.9923019
```

Be careful now... when you go to the reduced model, you no longer have the Epith and USize... so

```
> nd2 <- data.frame(Adhes=1, BNucl=1, Chrom = 3, Mitos = 1, NNucl=1, Thick=4,
UShap=1)
> predict(reduced, nd2, type="response")
1
0.9921115
```

Both the full model and the reduced model predict "Benign" for this person. In case you're interested, you can also get the predicted values like this:

```
> x0<-c(1,1,1,3,1,1,4,1)
> ilogit1<-sum(x0*coef(reduced))
> ilogit1
```

```
[1] 4.834428
> ilogit(ilogit1)
[1] 0.9921115
```

We can also get a 95% prediction interval for the fit like this:

```
> predict(reduced, newdata=data.frame(Adhes=1, BNucl=1, Chrom=3, Mitos=1,
NNucl=1, Thick=4, UShap=1), se=T)
$fit
      1
4.834428

$se.fit
[1] 0.5815185

$residual.scale
[1] 1
```

These numbers refer to the linear prediction... we can use them to make a 95% CI for the logistic regression $p()$ function at this set of predictor values like this:

```
> ilogit(c(4.834428-1.96*.5815185,4.834428+1.96*.5815185))
[1] 0.9757467 0.9974629
```

That is, we can be about 95% confident that, provided all the model assumptions are met, that the true $p()$ value for Adhes=1, BNucl=1, etc. is between 0.9757467 and 0.9974629.

- (d) Suppose a tumor is classified as benign if $p > 0.5$ and classified as malignant if $p < 0.5$ (remember “1” means benign and “0” indicates malignant for the Class variable). Compute the number of errors of both types that will be made if this method is applied to the current data with the reduced model. Also, compute the percentage of classifications that result in each kind of error (that is, compute the error rates- false positive and false negative).

```
> table1<-predict(reduced, type="response")
> table(wbca$Class, 1*table1 > .5)

FALSE TRUE
0  227  11
1   9 434
```

That is, the number of malignant tumors that were falsely categorized as benign was 9. The number of benign tumors falsely categorized as malignant was 11.

Accuracy = $(227 + 434) / 681 = .97063$

False positive rate = $11 / 238 = 0.0462$ (I'm assuming “positive” means malignant)

False negative rate = $9 / 443 = 0.0203$

- (e) Suppose we move the cut-off point to 0.9 so that $p < 0.9$ indicates malignant and $p > 0.9$ indicates benign. What are the new error counts and rates?

```
> table(wbca$Class, 1*table1 > .9)
```

```
FALSE TRUE
0  237   1
1   16 427
```

Accuracy = $(237 + 427) / 681 = 0.97503$

False positive rate = $1 / 238 = 0.004202$ (I'm assuming "positive" means malignant)

False negative rate = $16 / 443 = 0.036117$

The false positive rate has decreased, but the false negative rate has increased. I would prefer to see the false negative rate go down.

- (f) It can be misleading to use the same data to fit a model and test its predictive ability. So split the original dataset into two parts: assign every third record to the test set and all the others to the training set. Use the training set to determine a good model (repeat parts (a) and (b)). Then use the test set to assess predictive performance (repeat parts (d) and (e)).

```
> Obs3rd<-which((1:681)%%3 ==0)
> test<-wbca[Obs3rd,]
> training <- wbca[-Obs3rd,]
> modelTrain<-glm(Class~.,data=training, family=binomial(logit))
> summary(modelTrain)
```

Call:

```
glm(formula = Class ~ ., family = binomial(logit), data = training)
```

Deviance Residuals:

```
Min      1Q  Median      3Q      Max
-1.98138 -0.00954  0.03310  0.07084  3.07275
```

Coefficients:

```
Estimate Std. Error z value Pr(>|z|)
(Intercept) 12.0244    2.0462  5.876 4.19e-09 ***
Adhes       -0.4859    0.1555 -3.126 0.00177 **
BNucl       -0.3732    0.1292 -2.888 0.00388 **
Chrom       -0.6655    0.2536 -2.625 0.00868 **
Epith        0.1779    0.2148  0.828 0.40744
Mitos       -0.6075    0.5103 -1.190 0.23388
NNucl       -0.5168    0.1828 -2.828 0.00469 **
Thick       -0.6533    0.2044 -3.197 0.00139 **
UShap       -0.5291    0.2612 -2.026 0.04280 *
```

```
USize      0.2672  0.2320  1.152  0.24947
```

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 592.796 on 453 degrees of freedom  
Residual deviance: 57.651 on 444 degrees of freedom  
AIC: 77.651
```

Number of Fisher Scoring iterations: 9

```
> pchisq(deviance(modelTrain), df.residual(modelTrain), lower=FALSE)  
[1] 1
```

```
> redTrain<-step(modelTrain)
```

Start: AIC=77.65

```
Class ~ Adhes + BNucl + Chrom + Epith + Mitos + NNucl + Thick +  
      UShap + USize
```

	Df	Deviance	AIC
- Epith	1	58.340	76.340
- USize	1	58.880	76.880
<none>		57.651	77.651
- Mitos	1	60.712	78.712
- UShap	1	61.450	79.450
- Chrom	1	65.983	83.983
- BNucl	1	67.373	85.373
- NNucl	1	67.538	85.538
- Adhes	1	68.073	86.073
- Thick	1	71.162	89.162

Step: AIC=76.34

```
Class ~ Adhes + BNucl + Chrom + Mitos + NNucl + Thick + UShap +  
      USize
```

	Df	Deviance	AIC
- USize	1	59.536	75.536
<none>		58.340	76.340
- Mitos	1	61.264	77.264
- UShap	1	61.702	77.702
- Chrom	1	66.515	82.515
- BNucl	1	67.402	83.402
- NNucl	1	67.556	83.556
- Adhes	1	68.310	84.310
- Thick	1	72.311	88.311

Step: AIC=75.54

Class ~ Adhes + BNucl + Chrom + Mitos + NNucl + Thick + UShap

```
      Df Deviance   AIC
<none>      59.536 75.536
- UShap  1  61.894 75.894
- Mitos  1  62.329 76.329
- Chrom  1  66.762 80.762
- NNucl  1  67.576 81.576
- BNucl  1  68.332 82.332
- Adhes  1  68.359 82.359
- Thick  1  72.363 86.363
>
```

```
> pchisq(deviance(redTrain), df.residual(redTrain), lower=FALSE)
[1] 1
```

```
> summary(redTrain)
```

Call:

```
glm(formula = Class ~ Adhes + BNucl + Chrom + Mitos + NNucl +
    Thick + UShap, family = binomial(logit), data = training)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.03312	-0.01224	0.04042	0.08373	2.85056

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	11.5571	1.8285	6.321	2.6e-10 ***
Adhes	-0.4249	0.1441	-2.949	0.00318 **
BNucl	-0.3341	0.1187	-2.815	0.00487 **
Chrom	-0.5963	0.2422	-2.462	0.01382 *
Mitos	-0.5822	0.4872	-1.195	0.23207
NNucl	-0.4192	0.1604	-2.614	0.00895 **
Thick	-0.6037	0.1924	-3.138	0.00170 **
UShap	-0.2943	0.2034	-1.447	0.14795

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 592.796 on 453 degrees of freedom
Residual deviance: 59.536 on 446 degrees of freedom
AIC: 75.536

Number of Fisher Scoring iterations: 9

```
> table2 <- predict(redTrain, newdata = test, type = "response", se = T)
> table(test$Class, 1*(table2$fit > .9))
```

```
  0  1
0 73  2
1  3 149
```

Accuracy = $(73 + 149) / (73 + 2 + 3 + 149) = 0.977974$

False positive rate = $2/75 = 0.0266667$

False negative rate = $3/152 = 0.01974$

- (g) Discuss how you could search for the “best” cut-off point to use for classifying tumors. Then write an R program to carry it out. What is the “best” cut-off value? What do you mean by “best”?

First, you would essentially want to determine which of the classification rates is most important. In the case of cancer, I think the false negative rate is. So what one could do is build a bunch of models for varying values of the cut-off (say, 0.01, 0.02, 0.03, ..., 0.98, 0.99) on the training set. Then for each one, evaluate the false negative rate. Choose the value of the cut-off that minimizes the false negative rate.

Of course, the hospital administrators and insurance providers might be more concerned with the false positive rate, because if a test is positive, then maybe further testing (which is expensive) is required... They might arrive at a different model... or perhaps some weighting of the two error rates would be appropriate....